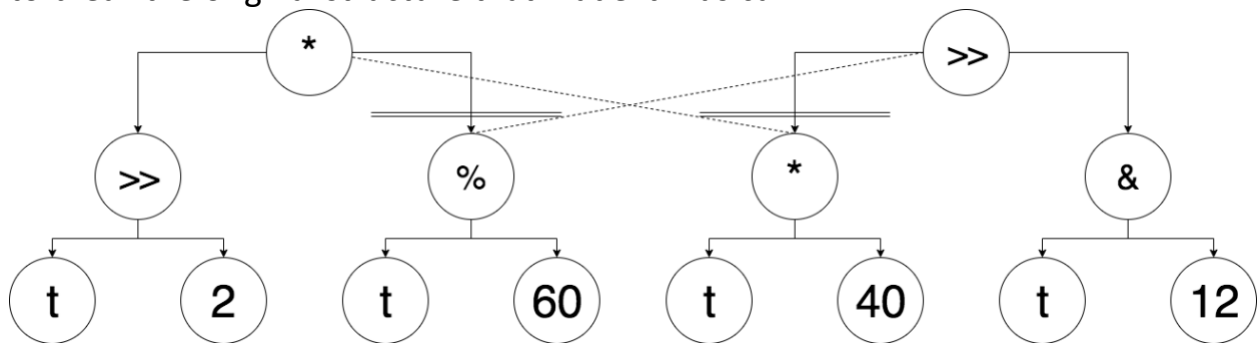


Progress Report 5: Crossover Strategy, Tree Feature Extraction & Support Vector Regression

This week I focused on finding a crossover strategy for our tree chromosomes, this allows us to combine trees from different clusters and mimic nature by “mating” them. I also spent some time extracting tree-based features in order to observe the relationship between or ByteBeat tracks and their underlying structure. Lastly, I attempted to use a Support Vector Regression to get a better prediction of ratings.

Starting with the crossover strategy, my approach is to select two trees from different sound-based clusters, and have them swap sub-trees. Sub-trees of lower heights have a higher probability of swapping while higher ones swap quite rarely. Additionally, the sub-trees are swapped at equal heights as shown in the figure below. The probability of crossover in our Genetic algorithm should be low, considering this is a significant mutation to the original tree and has the potential to break the original structure that made it musical.




Once the trees could freely mutate and mate, I wanted to analyze these them and how they relate to the music they generate. For this, I extracted the following tree-based features and ran a k-mean clustering algorithm based on them.

- Height: height of the tree.
- Leaves: number of leaves.
- t-count: number of "t" leaves.
- Operators: number of operators.
- Count of each operator: 9 columns counting each operator.
- Powers of 2: number of leaf nodes that are powers of 2.
- Average Operand: Average value of operands

The clustering revealed some interesting patterns. The highest rated group contained many bit-shift and bit-wise "&" operators, which is in-line with Viznut's

research on the Sierpinski Harmonies. The lowest rated group had a surprisingly high-amount of powers of two, which runs contrary to Viznut's hypothesis about these numbers.



| | >> | * | & | ^ | % | | - | + | / | 2_powers |
|---|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|
| 0 | 1.394366 | 0.926056 | 1.183099 | 0.647887 | 0.306338 | 0.461268 | 0.264085 | 0.281690 | 0.059859 | 0.669014 |
| 1 | 1.742331 | 1.153374 | 3.509202 | 0.828221 | 0.257669 | 0.644172 | 0.325153 | 0.282209 | 0.018405 | 0.822086 |
| 2 | 3.867470 | 1.271084 | 2.030120 | 0.668675 | 0.240964 | 0.759036 | 0.240964 | 0.259036 | 0.030120 | 1.162651 |
| 3 | 2.035088 | 1.175439 | 1.298246 | 0.842105 | 0.482456 | 2.701754 | 0.307018 | 0.333333 | 0.052632 | 1.236842 |
| 4 | 2.223214 | 1.401786 | 2.107143 | 1.625000 | 0.473214 | 1.000000 | 0.571429 | 0.491071 | 0.017857 | 3.312500 |
| 5 | 1.813665 | 3.316770 | 1.732919 | 0.919255 | 0.242236 | 0.708075 | 0.285714 | 0.273292 | 0.068323 | 1.298137 |

The last bit of progress I made this week was trying out different models to try to predict a song's rating based on its tree and spectral features. So far I have had little success in generating accurate predictions, but the best model so far has been the support-vector regression. I will meet with my Machine Learning TA to try to solve this issue.